

# **Process vs Thread Performance Analysis**

## **Comprehensive Benchmark Suite in C**

**Roll Number: 25081**

**Name: Saloni**

**Course: CSE638 - Graduate Systems**

**Project: PA01 - System Performance  
Analysis**

**Language: C with POSIX APIs**

**GitHub:**

**[https://github.com/saloninarang27/25081\\_PA01](https://github.com/saloninarang27/25081_PA01)**

# INDEX

SNO	TITLE	PAGE NO
1	Introduction	3
2	Experiment Set-up	3
3	Project Overview	3
4	Implementation Details (Program A + Program B)	3,4
5	AI Declaration	4
6	Part C and Part D tables	4-6
7	Screenshots and Plots	7-10
8	Analysis and Observations	11
9	Conclusion	11

## **Introduction**

Modern operating systems provide multiple concurrency mechanisms to improve performance and resource utilization. Two widely used mechanisms are processes and threads. While processes provide strong isolation, threads are lightweight and share the same address space. In this assignment, we experimentally evaluate the performance differences between process-based concurrency and thread-based concurrency under different workload types:

- CPU-bound
- Memory-bound
- I/O-bound

## **Experimental Setup**

### Hardware and OS

- Linux-based system
- Single-core pinning using taskset to avoid scheduler interference

### Measurement Tools

- top – CPU and memory usage
- iostat – disk I/O statistics
- time – execution duration

### Programs Implemented

- Program A: Uses fork() to create multiple processes
- Program B: Uses pthread\_create() to create multiple threads
- Each program executes one of the three worker functions – cpu, mem, io.

## **Project Overview**

The PA01 project implements a complete benchmarking suite consisting of four parts:

### **Part A: Basic Implementation**

- **Program A:** Multi-process implementation using fork()
- **Program B:** Multi-threaded implementation using pthread\_create()

### **Part B: Worker Functions**

- **CPU Worker:** 1B floating-point operations
- **Memory Worker:** 200MB array with cache-hostile access patterns
- **I/O Worker:** 10MB file operations with read/write verification

### **Part C & D: Benchmarking and Analysis**

- **Part C:** Baseline metrics with fixed 2 processes/threads
- **Part D:** Scaling analysis from 2 to 8 processes/threads with visualization

## **Implementation Details**

### **Concurrency Models**

#### **Program A: Process-Based (fork)**

- Parent process creates N child processes via fork()
- Each child independently executes worker function
- Parent waits for all children with waitpid()
- Strong isolation: separate address spaces, page tables, file descriptors
- Higher context switch cost due to address space switching

## Program B: Thread-Based (pthread)

- Main thread creates N worker threads via pthread\_create()
- All threads share same address space and file descriptors
- Main thread synchronizes with pthread\_join()
- Weak isolation: shared memory allows race conditions
- Lower context switch cost (same memory space)

## AI Declaration

### AI-GENERATED COMPONENTS:

AI Assistance Used For:

- Scripting logic and system tool integration
- Data visualization patterns
- Build system configuration
- Worker algorithm implementations
- Core API understanding (fork/pthread/wait)
- Documentation and explanation

### 3.1. Part C: Baseline Benchmarking Results

Baseline performance metrics at a scale of 4 workers:

Program+Worker	CPU%	Memory(KB)	IO	Time(s)
progA+cpu	98.40	10368	83.37	5.83
progA+mem	98.55	420608	125.16	22.56
progA+io	7.00	10752	138.51	21.14
progB+cpu	98.18	9088	69.27	6.06
progB+mem	98.53	418688	347.49	22.00
progB+io	5.72	9216	139.83	21.22

### 3.2.1. Part D: CPU Worker Scaling Results

Program	Scale	Avgcpu Percent	Avgmemory Kb	Totalio Kb	Executiontime Sec
progA	2	72.75	10368	5.74	4.24
progA	3	97.84	11008	17.73	6.39
progA	4	95.98	11520	5.7	8.44
progA	5	101.11	12288	5.7	10.69
progB	2	75.00	8960	9.66	4.25
progB	3	100.00	8960	5.66	6.40
progB	4	100.00	9088	5.63	9.30
progB	5	97.97	9088	377.61	10.59
progB	6	90.90	9088	5.62	12.71
progB	7	91.67	9088	5.59	15.19
progB	8	99.41	9088	5.56	17.23


### 3.2.2. Part D: MEM Worker Scaling Results


<b>Program</b>	<b>Scale</b>	<b>Avgcpu Percent</b>	<b>Avgmemory Kb</b>	<b>Totalio Kb</b>	<b>Executiontime Sec</b>
progA	2	91.49	420480	177.73	18.85
progA	3	96.55	626176	13.73	27.05
progA	4	102.31	831872	161.7	36.64
progA	5	99.34	1037568	181.69	46.68
progB	2	92.61	418560	25.66	18.89
progB	3	96.69	623488	69.66	27.10
progB	4	95.89	828288	169.63	37.21
progB	5	98.74	7425	81.65	45.22
progB	6	99.85	7425.2	205.61	54.63
progB	7	99.04	9404	357.58	66.52
progB	8	97.12	7425.6	245.55	73.75

### 3.2.3. Part D: IO Worker Scaling Results


<b>Program</b>	<b>Scale</b>	<b>Avgcpu Percent</b>	<b>Avgmemory Kb</b>	<b>Totalio Kb</b>	<b>Executiontime Sec</b>
progA	2	6.00	10624	5.74	11.60
progA	3	7.00	11392	61.7	11.91
progA	4	9.23	12160	169.69	15.87
progA	5	5.00	12928	33.68	19.82
progB	2	3.23	9088	157.65	10.93
progB	3	7.00	9088	5.64	11.84
progB	4	5.00	9088	17.62	16.40
progB	5	6.36	9088	41.62	16.86
progB	6	9.88	9088	137.59	19.91
progB	7	8.18	9088	45.58	27.53
progB	8	9.42	9088	141.53	27.76

## Screenshots And Plots






## Execution Time Comparison (All Worker Types)



**Execution Time - MEM Worker**



**Execution Time - IO Worker**



```

saloni@Saloni: ~
top - 17:09:00 up 13 min, 1 user, load average: 0.45, 0.10, 0.03
Tasks: 30 total, 3 running, 27 sleeping, 0 stopped, 0 zombie
%Cpu(s): 25.0 us, 0.1 sy, 0.0 ni, 74.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 7816.2 total, 7360.1 free, 461.4 used, 143.6 buff/cache
MiB Swap: 2048.0 total, 2048.0 free, 0.0 used. 7354.8 avail Mem

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
733 saloni 20 0 3612 768 768 R 98.7 0.0 0:14.49 program_a
734 saloni 20 0 3612 768 768 R 98.7 0.0 0:14.49 program_a
1 root 20 0 21528 12232 9288 S 0.0 0.2 0:00.63 systemd
2 root 20 0 3060 1792 1792 S 0.0 0.0 0:00.01 init-systemd(Ub
6 root 20 0 3060 1792 1792 S 0.0 0.0 0:00.00 init
42 root 19 -1 66820 18712 17944 S 0.0 0.2 0:00.20 systemd-journal
89 root 20 0 25136 6272 4992 S 0.0 0.1 0:00.14 systemd-udevd
144 systemd+ 20 0 21456 12672 10624 S 0.0 0.2 0:00.14 systemd-resolve

```

```

saloni@Saloni:/mnt/c/Users/safet/Desktop/New/25081_PA01$ bash MT25081_Part_C_benchmark.sh
Checking for required tools...
All required tools found.

```

PA01 PART C: BASELINE BENCHMARKING – PROCESSES VS THREADS

Roll Number: 25081

Scale: 2 workers (pinned to a single CPU core)

Objective: Establish baseline metrics with a fixed scale (2) to compare single-core process vs. thread performance.

Initializing CSV file: MT25081\_Part\_C\_CSV.csv  
CSV initialized with headers

System Information:

CPU Cores: 8

Start Time: 2026-01-23 15:42:17

Running 6 baseline benchmark combinations...

Running: progA+cpu

```

DEBUG: Started /mnt/c/Users/safet/Desktop/New/25081_PA01/progA (cpu) with PID: 441
[progA] Starting 2 processes with worker type: cpu
[progA] Parent waiting for 2 children to finish...
[progA] Child process 1 (PID: 444) started

```

```

DEBUG: Parsed CPU: 91, Parsed MEM: 10368
[progA] Child process 2 (PID: 1771) completed
[progA] Child process 1 (PID: 1770) completed
[progA] Child 1 exited with status: 0
[progA] Child 2 exited with status: 0
[progA] All 2 children completed. Parent exiting.
DEBUG: Program with PID 1767 terminated.

```

Completed: progA+cpu

Avg CPU: 98.40%

Max Memory: 10368 KB

Total I/O Writes: 83.37 KB

Execution Time: 5.83s

```

DEBUG: Parsed CPU: 100, Parsed MEM: 420608
[progA] Child process 2 (PID: 1858) completed
[progA] Child process 1 (PID: 1857) completed
[progA] Child 1 exited with status: 0
[progA] Child 2 exited with status: 0
[progA] All 2 children completed. Parent exiting.
DEBUG: Program with PID 1854 terminated.

```

Completed: progA+mem

Avg CPU: 98.55%

Max Memory: 420608 KB

Total I/O Writes: 125.16 KB

Execution Time: 22.56s

```
[progA] Child process 1 (PID: 2126) completed
[progA] Child 1 exited with status: 0
[progA] Child process 2 (PID: 2127) completed
[progA] Child 2 exited with status: 0
[progA] All 2 children completed. Parent exiting.
DEBUG: Program with PID 2123 terminated.
Completed: progA+io
Avg CPU: 7.00%
Max Memory: 10752 KB
Total I/O Writes: 138.51 KB
Execution Time: 21.14s
```

```
-----+
DEBUG: Parsed CPU: 90.9, Parsed MEM: 9088
[progB] Thread 2 (TID: 133723724965568) completed
[progB] Thread 1 (TID: 133723733358272) completed
[progB] Thread 1 joined successfully
[progB] Thread 2 joined successfully
[progB] All 2 threads completed. Main thread exiting.
DEBUG: Program with PID 2379 terminated.
Completed: progB+cpu
Avg CPU: 98.18%
Max Memory: 9088 KB
Total I/O Writes: 69.27 KB
Execution Time: 6.06s
```

```
DEBUG: Parsed CPU: 90, Parsed MEM: 418688
[progB] Thread 1 (TID: 135004428433088) completed
[progB] Thread 1 joined successfully
[progB] Thread 2 (TID: 135004420040384) completed
[progB] Thread 2 joined successfully
[progB] All 2 threads completed. Main thread exiting.
DEBUG: Program with PID 2469 terminated.
Completed: progB+mem
Avg CPU: 98.53%
Max Memory: 418688 KB
Total I/O Writes: 347.49 KB
Execution Time: 22.00s
```

```
DEBUG: Parsed CPU: 9.1, Parsed MEM: 9216
[progB] Thread 2 (TID: 124871459727040) completed
[progB] Thread 1 (TID: 124871468119744) completed
[progB] Thread 1 joined successfully
[progB] Thread 2 joined successfully
[progB] All 2 threads completed. Main thread exiting.
DEBUG: Program with PID 2742 terminated.
Completed: progB+io
Avg CPU: 5.72%
Max Memory: 9216 KB
Total I/O Writes: 139.83 KB
Execution Time: 21.22s
```

```
✓ All baseline benchmarks completed successfully!
```

```
Results saved to: MT25081_Part_C_CSV.csv
```

```
CSV Contents:
```

```
-----+
Program+Worker,CPU%,Memory(KB),IO,Time(s)
progA+cpu,98.40,10368,83.37,5.83
progA+mem,98.55,420608,125.16,22.56
progA+io,7.00,10752,138.51,21.14
progB+cpu,98.18,9088,69.27,6.06
progB+mem,98.53,418688,347.49,22.00
progB+io,5.72,9216,139.83,21.22
-----+
```

## **Analysis and Observations**

### **CPU-Bound Workload**

- Both processes and threads achieve close to full CPU utilization (~98-100%), indicating effective core saturation
- Execution time is slightly lower for processes
- Memory usage is comparable in both models

**Analysis:** Threads complete CPU-intensive tasks faster than processes due to lower management overhead while achieving similar CPU utilization.

### **Memory-Bound Workload**

- Memory usage is approx. same for both processes and threads, with threads showing marginally lower memory consumption
- CPU utilization is high for processes but significantly lower for threads
- Thread-based execution completes faster than process-based execution

**Analysis:** Threads are more efficient in memory-bound workloads as shared address space reduces CPU overhead and execution time despite similar memory usage.

### **I/O Bound Workload**

- CPU utilization is very low for both models, confirming that execution is dominated by disk operations
- Execution time is significantly higher than CPU workloads
- The process-based implementation exhibits higher total disk I/O compared to the thread-based version

**Analysis:** I/O-bound workloads are dominated by disk latency, resulting in low CPU utilization and minimal performance difference between processes and threads.

## **Conclusion**

- Threads outperform processes for due to lower overhead and shared address space
- Memory-bound workloads show limited scalability, as memory bandwidth becomes the bottleneck
- I/O-bound workloads are dominated by disk performance, making the choice of concurrency model less significant.
- The assignment demonstrates that the choice between processes and threads must be guided by workload characteristics.
- Threads use significantly less memory for non-memory tasks because threads share an address space.